

Churn_model.py: Criar um **modelo de classificação** que estime a probabilidade de um cliente **dar churn** (parar de comprar / ficar inativo) num horizonte definido

```
# IA_e_ML/churn_model.py

import pandas as pd

import numpy as np

from sklearn.model_selection import train_test_split, GridSearchCV, StratifiedKFold

from sklearn.ensemble import RandomForestClassifier

from sklearn.pipeline import Pipeline

from sklearn.preprocessing import OneHotEncoder, StandardScaler

from sklearn.impute import SimpleImputer

from sklearn.compose import ColumnTransformer

from sklearn.metrics import roc_auc_score, classification_report, confusion_matrix

import joblib

import matplotlib.pyplot as plt


# carregue features (pode usar segmentos e rfm)

df = pd.read_csv("IA_e_ML/segmentos_clientes.csv") # já contém RFM e cluster

# Suponha que tenhamos target 'churn_90d' (1 = churn nos próximos 90 dias)

# Se não tem, você precisa construir baseado em last_active_date

if 'churn_90d' not in df.columns:

    # exemplo simples: churn se última compra > 90 dias

    df['churn_90d'] = (df['recency'] > 90).astype(int)


# features e target

target = 'churn_90d'

cat_cols = ['gender', 'region', 'cluster'] if 'gender' in df.columns else ['cluster']

num_cols = ['recency', 'frequency', 'monetary', 'age'] if 'age' in df.columns else

['recency', 'frequency', 'monetary']

X = df[num_cols + cat_cols]

y = df[target]
```

```

# split
X_train, X_test, y_train, y_test = train_test_split(X, y, stratify=y, test_size=0.2,
random_state=42)

# preprocessing
num_pipe = Pipeline([
    ('imputer', SimpleImputer(strategy='median')),
    ('scaler', StandardScaler())
])
cat_pipe = Pipeline([
    ('imputer', SimpleImputer(strategy='most_frequent')),
    ('ohe', OneHotEncoder(handle_unknown='ignore', sparse=False))
])
preproc = ColumnTransformer([
    ('num', num_pipe, num_cols),
    ('cat', cat_pipe, cat_cols)
])

pipe = Pipeline([
    ('preproc', preproc),
    ('clf', RandomForestClassifier(n_estimators=200, random_state=42, n_jobs=-1))
])

# hyperparam tuning (pequeno)
param_grid = {
    'clf__max_depth': [5, 10, None],
    'clf__min_samples_split': [2, 5]
}
cv = StratifiedKFold(n_splits=4, shuffle=True, random_state=42)
gs = GridSearchCV(pipe, param_grid, scoring='roc_auc', cv=cv, n_jobs=-1)
gs.fit(X_train, y_train)

```

```

# avaliar

best = gs.best_estimator_
y_pred_proba = best.predict_proba(X_test)[:,-1]
y_pred = best.predict(X_test)

auc = roc_auc_score(y_test, y_pred_proba)
report = classification_report(y_test, y_pred, output_dict=False)
cm = confusion_matrix(y_test, y_pred)

# salvar

joblib.dump(best, "IA_e_ML/churn_model.joblib")

with open("IA_e_ML/churn_metrics.txt", "w") as f:
    f.write(f"AUC: {auc:.4f}\n")
    f.write(report)
    f.write("\nConfusion matrix:\n")
    f.write(str(cm))

# plot ROC (simples)

from sklearn.metrics import roc_curve
fpr, tpr, _ = roc_curve(y_test, y_pred_proba)
plt.figure()
plt.plot(fpr, tpr)
plt.plot([0,1],[0,1], '--')
plt.xlabel("FPR"); plt.ylabel("TPR"); plt.title(f"ROC curve (AUC={auc:.3f})")
plt.savefig("IA_e_ML/roc_curve.png", dpi=150)

```